

High-Dimensional Statistical Modeling and Analysis Of Custom Integrated Circuits

Trent McConaghy

Co-founder & CSO, Solido Design Automation

Custom Integrated Circuits Conference (CICC)

San Jose, CA, Sept 2011





furry robot



Search Images

- Everything
- Images
- Videos
- News
- More

Sort by relevance

Sort by subject

Any size

Large

Medium

Icon

Larger than...

Exactly...

Any color

Full color

Black and white



Any type

Face

Photo

Clip art







Line drawing

Standard view

Show sizes

Any time

Past week



Elmo's robot
the gadgets us
270 x 349 - ...
Roboaplan R3 Media robot In
Similar More news



- Everything
- Images
- Videos
- News
- More

Sort by relevance
Sort by subject

Any size

Large
Medium
Icon
Larger than...
Exactly...


Any color

Full color
Black and white



Any type

Face
Photo
Clip art
Line drawing

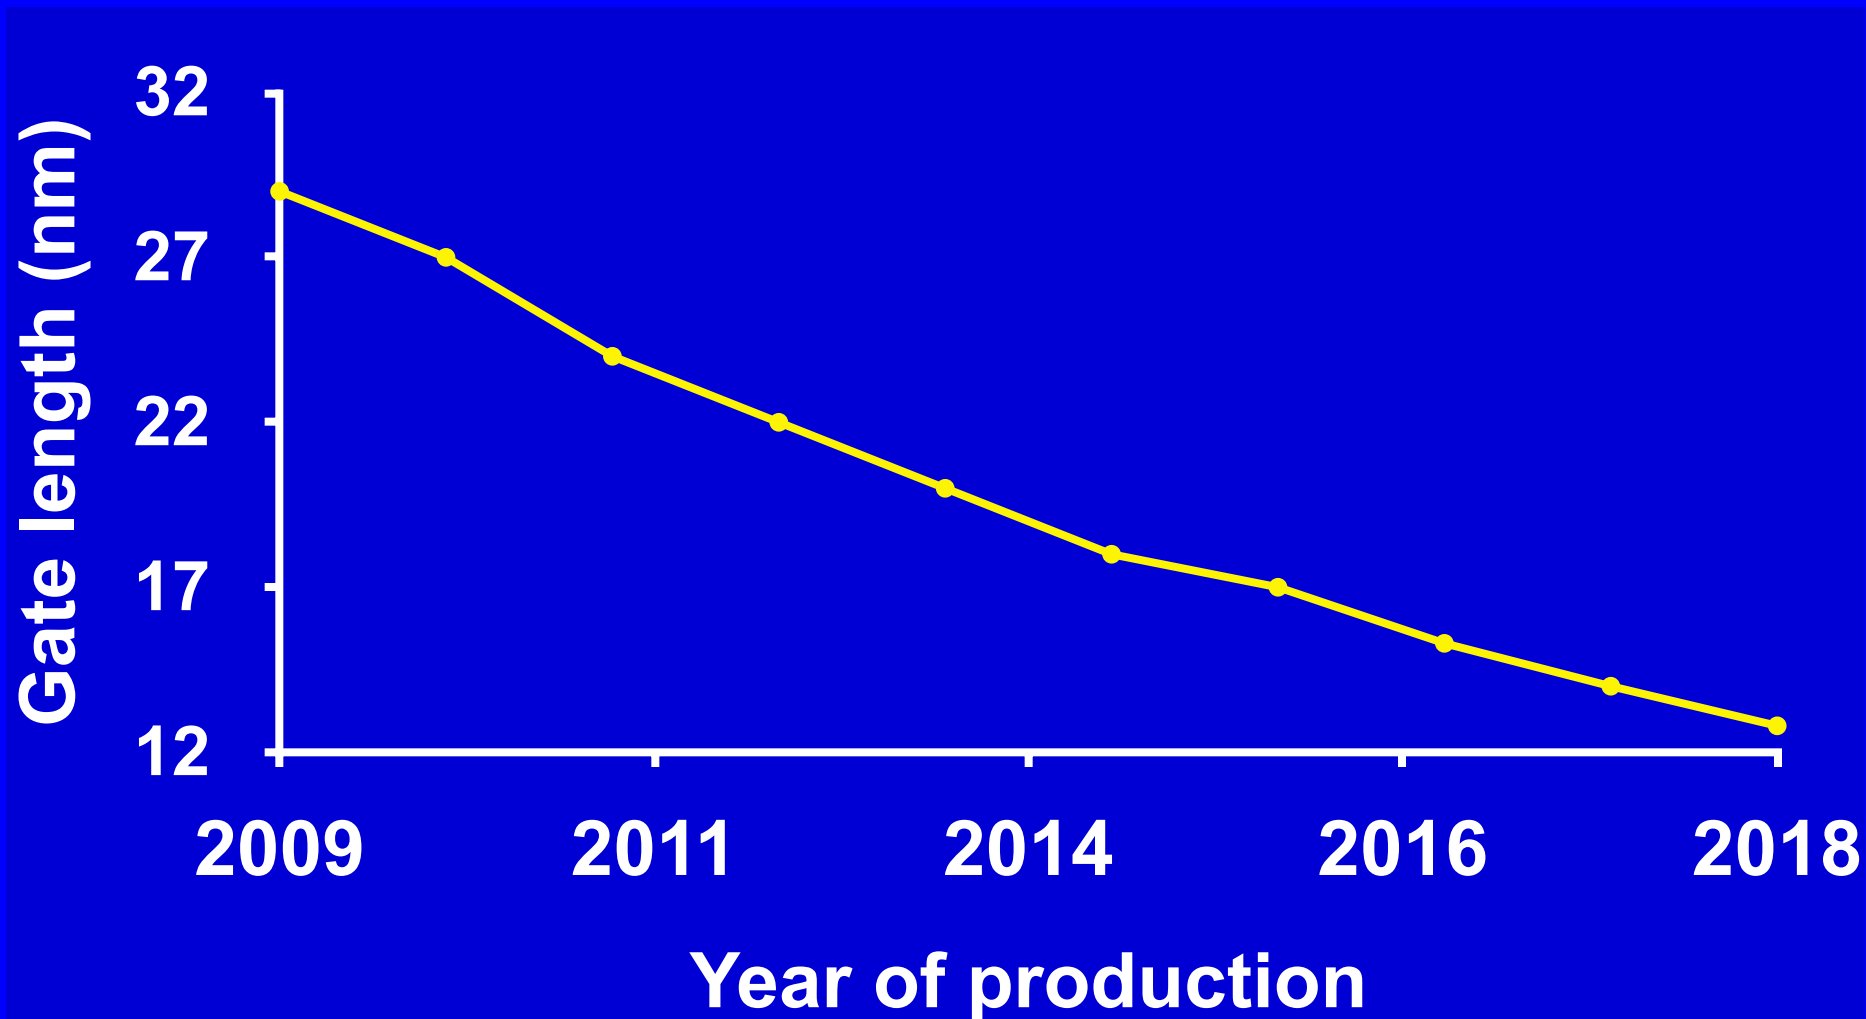


How *does* Google find furry robots?
(Not the aim of this talk, but we'll find out...)


Outline

- Motivation
- Proposed flow
- Background
- Fast Function Extraction (FFX)
- Results
- Conclusion


With Moore's Law, Transistors Are Shrinking...




Transistors are shrinking But atoms aren't!



traditional



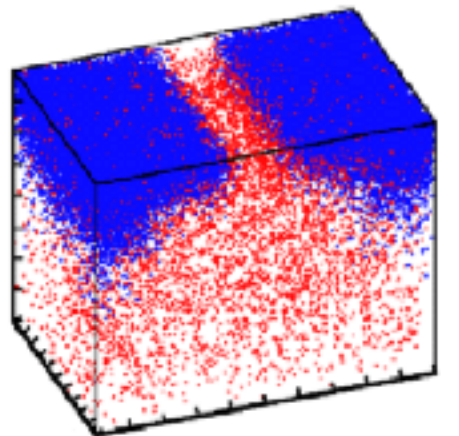
22nm




sub 10 nm

Variation = atoms out of place


Relative variation worsens as transistors shrink



Random dopant effects



Oxide thickness

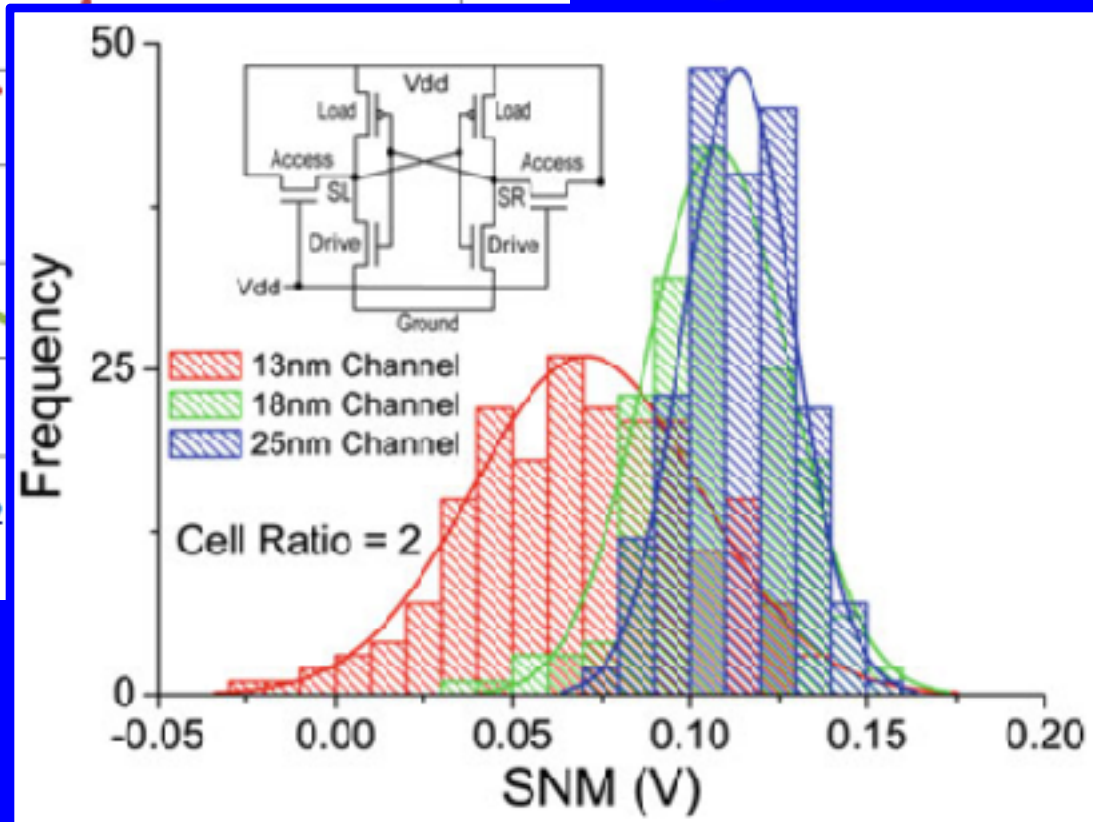



- A. Asenov, "Statistical Nano CMOS Variability and Its Impact on SRAM", Extreme Statistics, Springer, 2010
- C. Visweswariah, "Mathematics and Engineering: A Clash of Cultures?", IBM, 2005
- ITRS, 2006


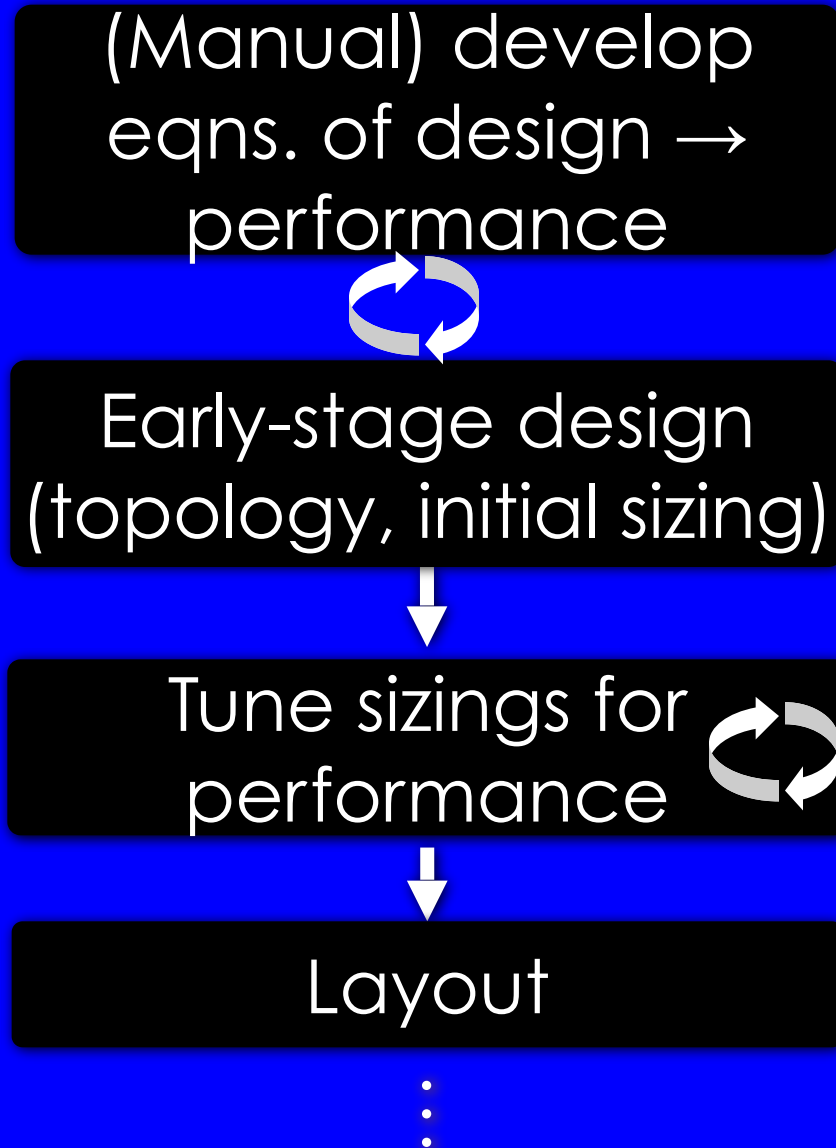
Higher device variation

→ Higher performance variation


→ Lower yield



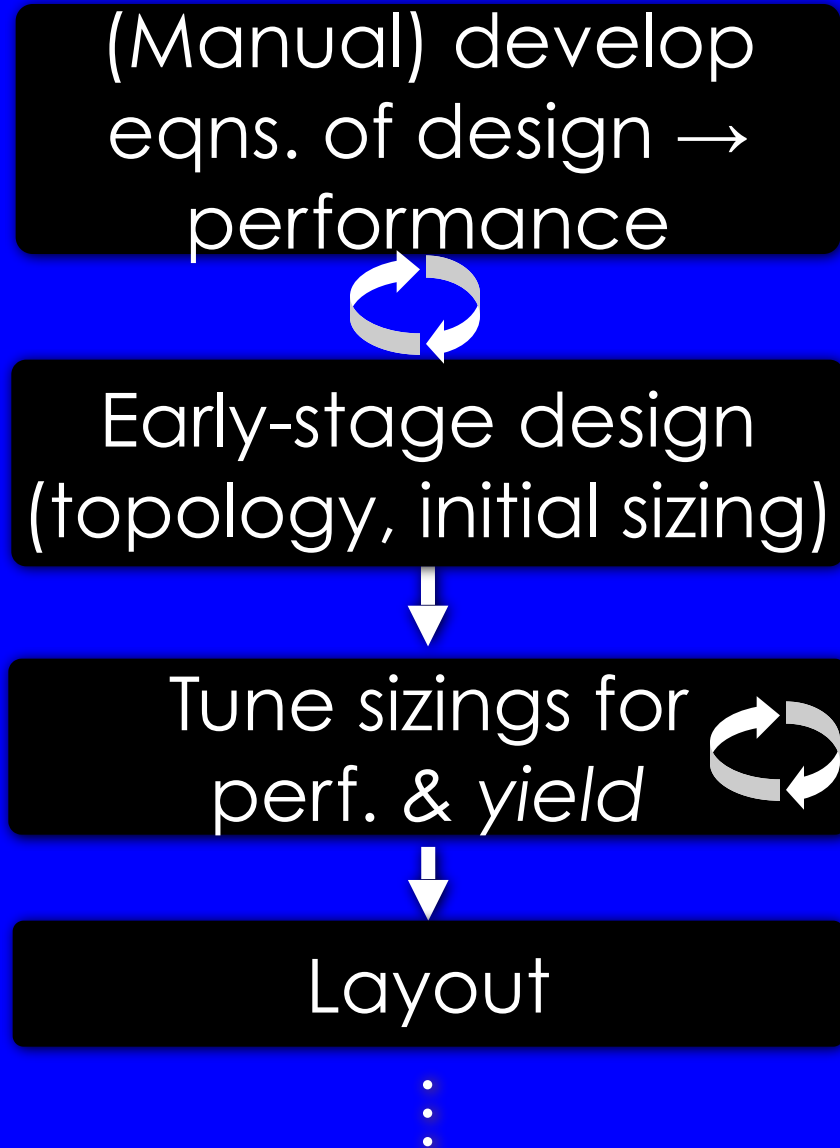
Typical Custom Design Flow



$W1, L1, W2, L2, ..$



Variation-Aware Design Flow




Variation-Aware Approaches:

- Direct Monte Carlo
- Design-specific 3σ corners
- ...

Direct Monte Carlo (MC) For Variation-Aware Design


(Manual) develop
eqns. of design \rightarrow
performance



Early-stage design
(topology, initial sizing)




Tune sizings for
perf. & yield
via direct MC



Layout

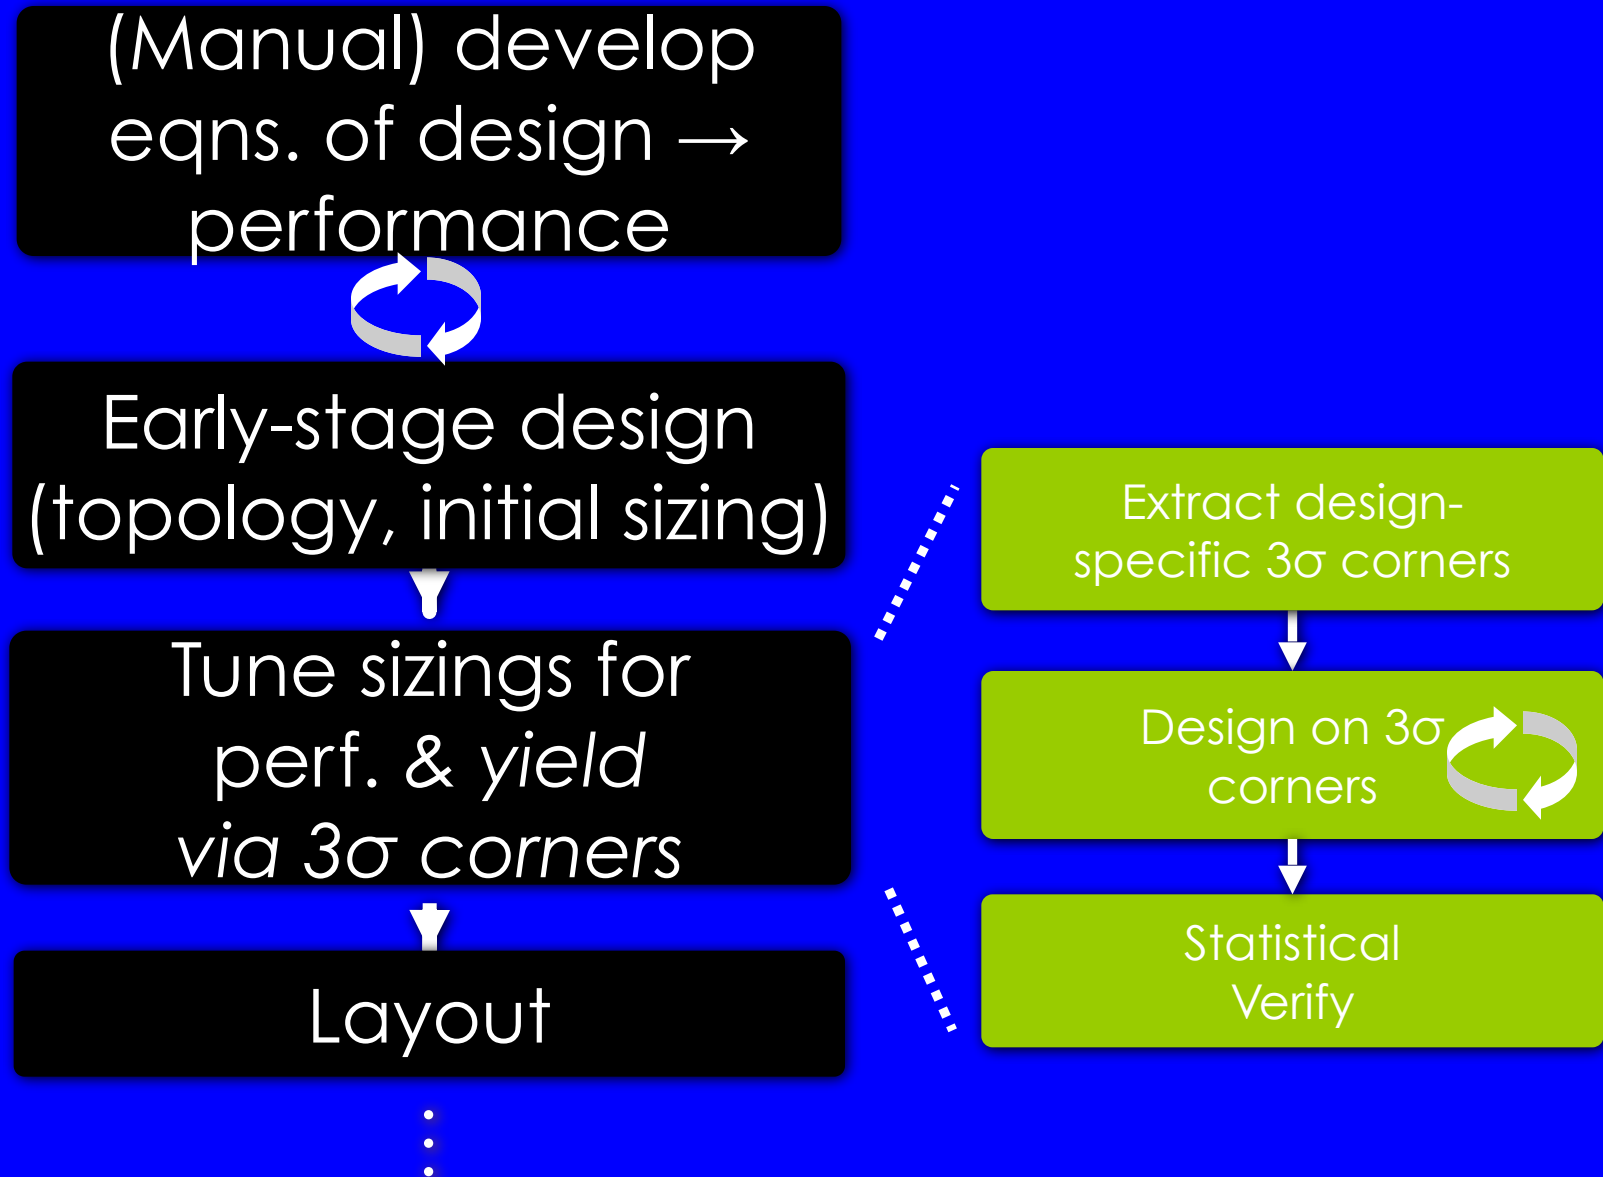
⋮

SPICE on all 50
(or 1K!) (or 10K!)
MC samples


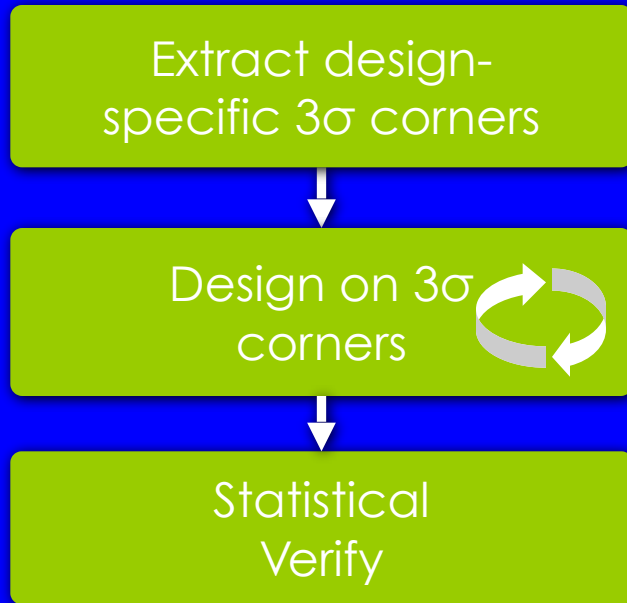


*Q: How to get MC out of the design loop,
yet retain accuracy of MC?*

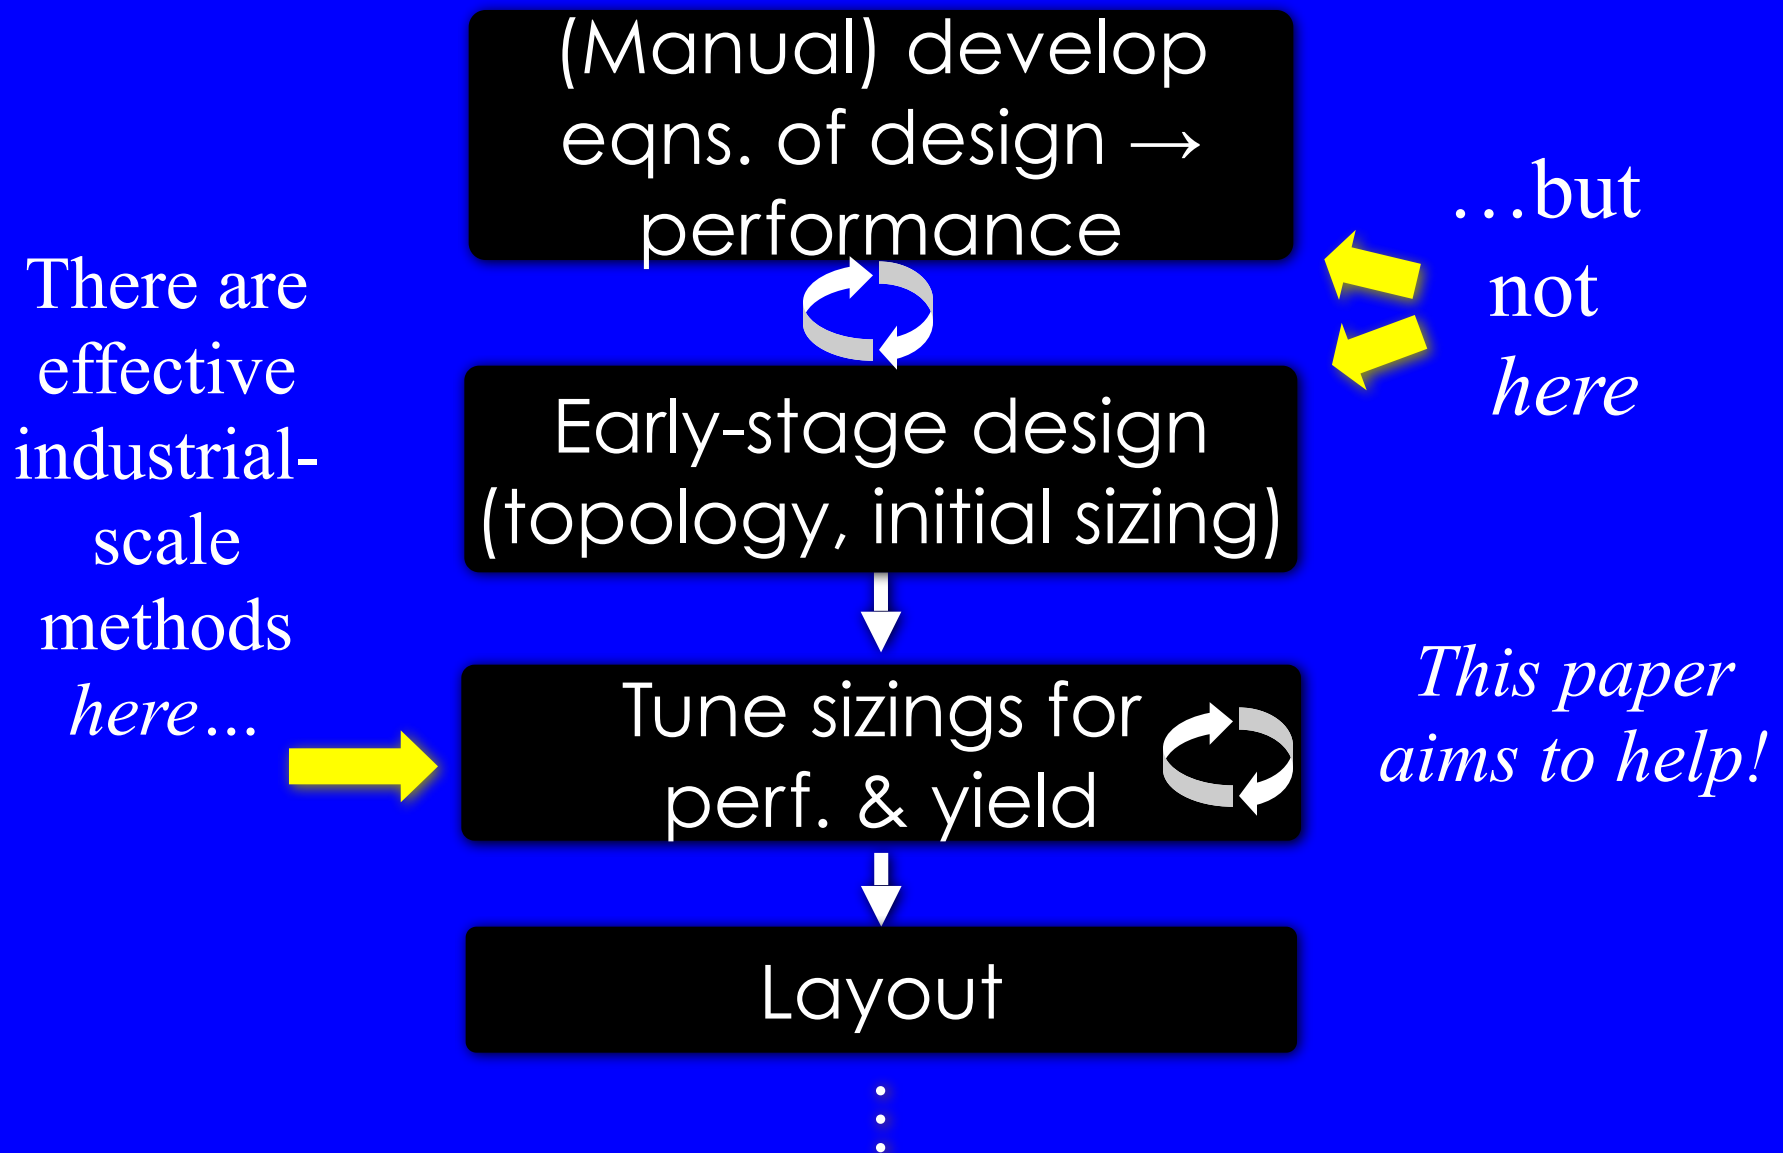
Design-Specific 3σ Corners For Variation-Aware Design



Design-Specific 3σ Corners For Variation-Aware Design



Variation-Aware Design Flow




Outline

- Motivation
- Proposed flow
- Backgrounder
- Fast Function Extraction (FFX)
- Results
- Conclusion

An Ideal Flow ... Is Infeasible

(Manual) develop eqns.
of design & *process vars.*
→ circuit performance



Early-stage design
(topology, initial sizing)



Tune sizings for
perf. & yield



Layout

⋮

“Dear designer,
To handle mismatch in early design
stages,

please add 1000 extra variables

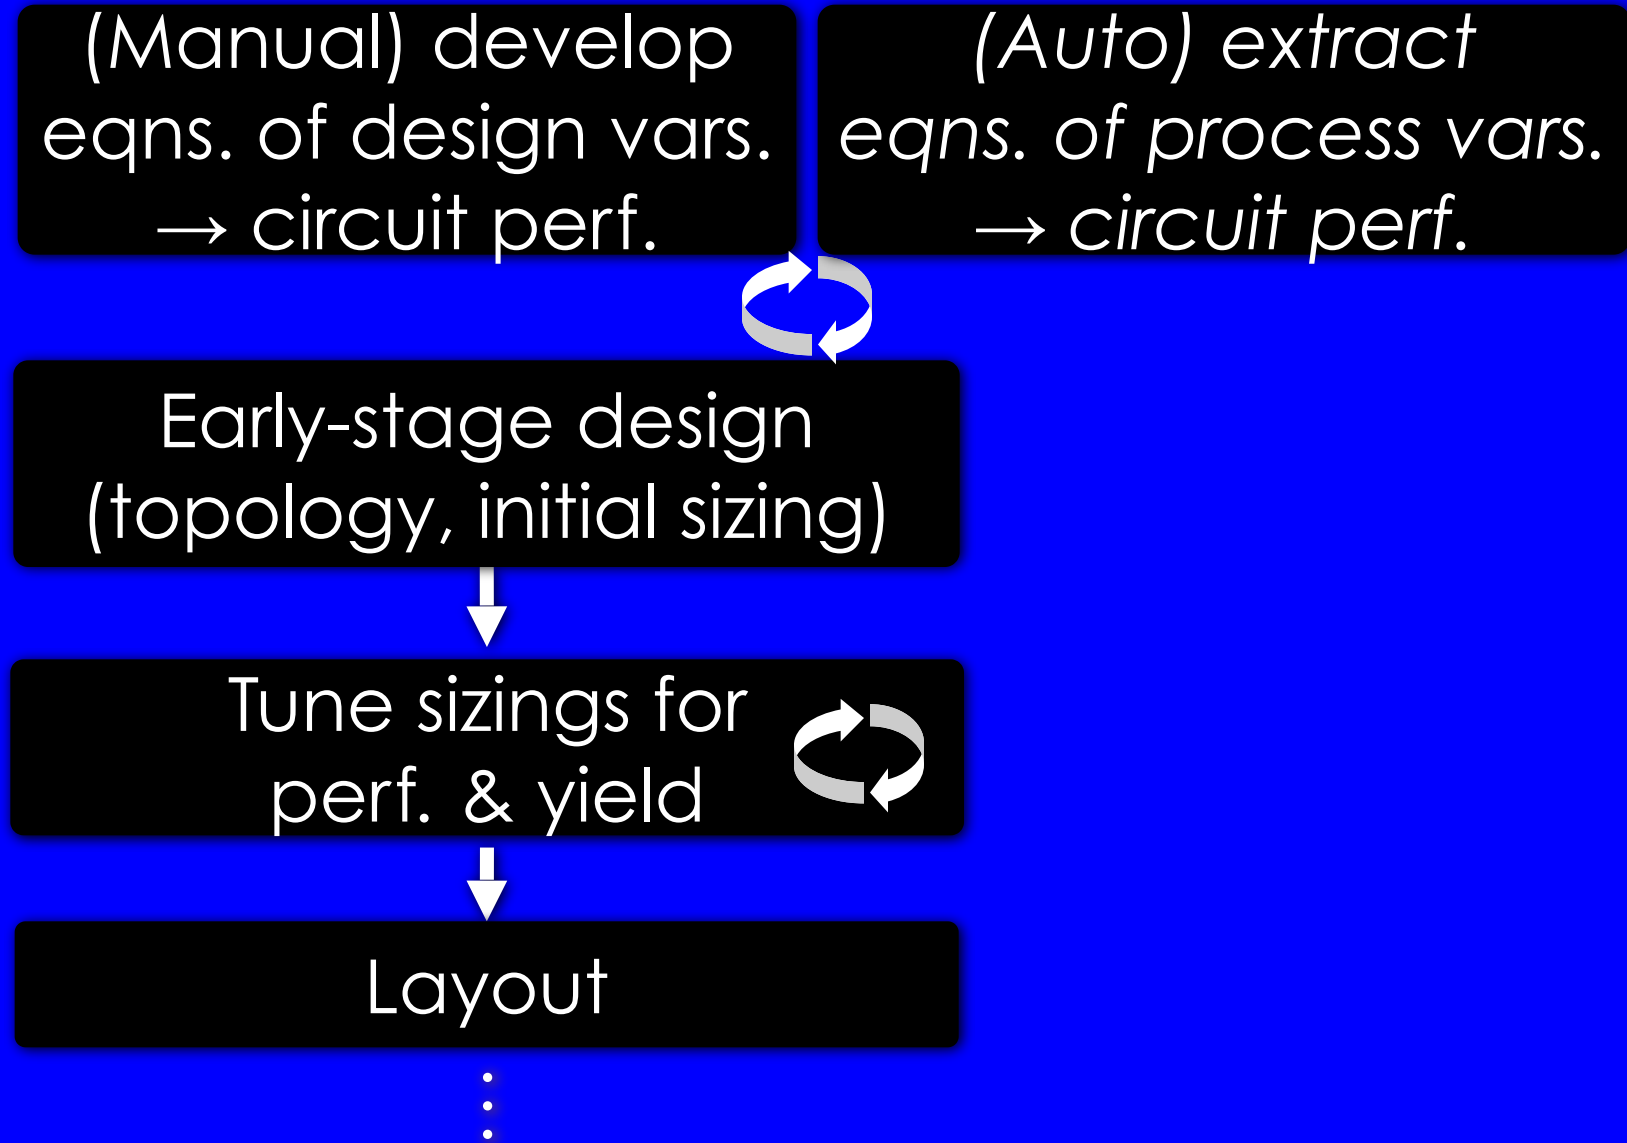
$\Delta\text{tox}_1, \Delta\text{Nsub}_1, \dots$

$\Delta\text{tox}_{100}, \Delta\text{Nsub}_{100}, \dots$ ”

Are *you* an expert at process
modeling?

Who is an expert at *both* process
modeling & topology development?

Proposed Flow




Equation-Extraction Step: Details, by Example (1D)

Draw 100-5000 MC
samples (X)

SPICE-Simulate on
samples (y)

Build whitebox
Model of $X \rightarrow y$



$$AV = 50.2 + 9.1 \cdot \Delta\text{tox}_1 + 3.2 \cdot \max(0, \Delta\text{tox}_1^2)$$

Equation-Extraction Step:

Problem Scope

- 5-10+ process variables per device
(e.g. BPV model)
- 10-100+ devices
- Therefore 50-1000+ input variables
- 100-5000 simulations (runtime cost)
- Need whitebox model, for insight!
- Ideally get a tradeoff of complexity vs. accuracy

Equation Extraction Step: Off-the-Shelf Modeling Options

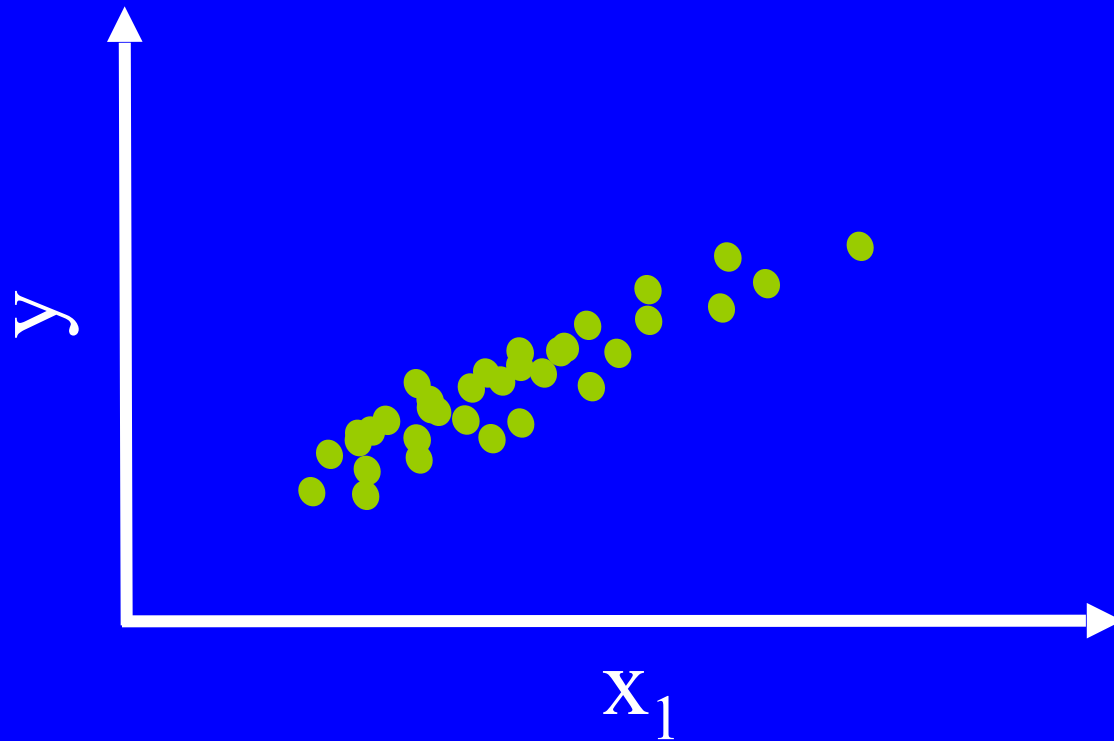
- Linear model – can't handle nonlinear
- Quadratic – not nonlinear enough, scaling?
- MARS, SVM, neural net, others – not whitebox

Nothing meets problem scope!

Outline


- Motivation
- Proposed flow
- Background
 - Least-squares regression
 - Regularized learning
- Fast Function Extraction (FFX)
- Results
- Conclusion

Given some x, y data...




...we can build 1D linear models

Many possible linear models!
(∞ to be precise)



1D Linear Least-Squares (LS) Regression

*Find linear model that minimizes $\sum (\hat{y}_i - y_i)^2$
(across all i in training data)*




1D Linear LS Regression

Find linear model that minimizes $\sum(\hat{y}_i - y_i)^2$

That is: $[w_0, w_1]^* = \text{minimize } \sum(\hat{y}_i - y_i)^2$

where $\hat{y}(x_1) = w_0 + w_1 \cdot x_1$



Solving this amounts to matrix inversion:


$$w = [1 \ X] / y$$

1D *Quadratic* LS Regression

$$[w_0, w_1, w_{11}]^* = \text{minimize } \sum (\hat{y}_i - y_i)^2$$

$$\text{where } \hat{y}(x_1) = w_0 + w_1 \cdot x_1 + w_{11} \cdot x_1^2$$

We are applying linear (LS) learning on linear & nonlinear basis functions. OK!




1D *Nonlinear* LS Regression

$$[w_0, w_1, w_{\sin}]^* = \text{minimize } \sum (\hat{y}_i - y_i)^2$$

$$\text{where } \hat{y}(x_1) = w_0 + w_1 \cdot x_1 + w_{\sin} \cdot \sin(x_1)$$


We are applying linear (LS) learning on linear & nonlinear basis functions. OK!



2D Linear LS Regression

$$[w_0, w_1, w_2]^* = \text{minimize } \sum (\hat{y}_i - y_i)^2$$


where $\hat{y}(x) = w_0 + w_1 \cdot x_1 + w_2 \cdot x_2$



2D Quadratic LS Regression

$$[w_0, w_1, w_2, w_{11}, w_{22}, w_{12}]^* = \text{minimize } \sum (\hat{y}_i - y_i)^2$$

$$\text{where } \hat{y}(x) = w_0 + w_1 \cdot x_1 + w_{11} \cdot x_1^2 + w_{22} \cdot x_2^2 + w_{12} \cdot x_1 \cdot x_2$$



Outline

- Motivation
- Proposed flow
- Background
 - Least-squares regression
 - Regularized learning
- Fast Function Extraction (FFX)
- Results
- Conclusion

Regularized Linear Regression

- Minimizes a combination of training error *and* *model sensitivity*

- Formulation:

$$w^* = \text{minimize} \left(\underbrace{\sum (\hat{y}_i(w) - y_i)^2}_{\text{minimize training error (fit training data better)}} + \underbrace{\lambda \cdot \sum |w_j|}_{\text{model sensitivity (reduce risk in future predictions)}} \right)$$

minimize training error (fit training data better)

model sensitivity
(reduce risk
in future predictions)

Regularized Linear Regression

When solving $w^* = \text{minimize } (\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|)$,
consider different values for λ ...



~~$\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|$~~ $\nearrow 0$

$w^* = [0, 0, \dots, 0]$
(constant response)



$\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|$ ~~$\nearrow 0$~~

w^* reduces to LS
solution


Pathwise Regularized Linear Regression



solve $w^* = \text{minimize } (\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|)$


at $\lambda = 1e40$ ($\lambda \rightarrow \infty$)

Pathwise Regularized Linear Regression




solve $w^* = \text{minimize} (\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|)$
at $\lambda = 1e30$

Pathwise Regularized Linear Regression




solve $w^* = \text{minimize} (\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|)$
at $\lambda = 1e20$

Pathwise Regularized Linear Regression




solve $w^* = \text{minimize} (\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|)$
at $\lambda = 1e10$

Pathwise Regularized Linear Regression




solve $w^* = \text{minimize} (\sum (\hat{y}_i(w) - y_i)^2 + \lambda \cdot \sum |w_j|)$
at $\lambda = 1e0$


Pathwise Regularized Linear Regression




Pathwise Regularized Linear Regression



Pathwise Regularized Linear Regression




Pathwise Regularized Linear Regression




Each column vector of w^* is a different linear model.
Which model is better / best? Accuracy vs. complexity...



Pathwise Regression Accuracy Per Model



Pathwise Regression Complexity Per Model



Accuracy – Complexity Tradeoff



Cool Properties of Pathwise Regression (versus plain old LS)

- Cool property #1: solving a full regularized path gives us accuracy-vs-complexity tradeoffs!
- Cool property #2: can have more coefficients than samples!
 - Regularization term $\sum |w_i|$ means no underdetermined problems
- Cool property #3: solving a regularized learning problem is just as fast (or faster) than solving a least-squares learning problem!
 - Why: convex optimization problem – one big hill

Q: How *does* Google find furry robots?

Answer (NIPS 2010):

1. Treat images as $1000 \times 1000 = 1\text{M}$ input vars. (x)
2. Crawl the web:
 - Find all images with “furry” & “robot” in filename. Assign y-value = 1
 - Find $\leq 5\text{K}$ images with “furry”, “robot”. y-value = 0.75
 - Randomly grab 5K more images. y-value = 0.0
3. Run regularized linear learning on $X \rightarrow y$ (learning on 1M input variables!!)
4. On all unseen images, run model. Return images sorted from highest \hat{y} down.



- Cool property #4: regularized learning can handle a *ridiculous* number of input variables!

Outline

- Motivation
- Proposed flow
- Backgrounder
 - Least-squares regression
 - Regularized learning
- Fast Function Extraction (FFX)
- Results
- Conclusion

Recap: Proposed Flow

(Manual) develop
eqns. of design vars.
→ circuit perf.

(Auto) extract
eqns. of process vars.
→ circuit perf.

Early-stage design
(topology, initial sizing)

Tune sizings for
perf. & yield

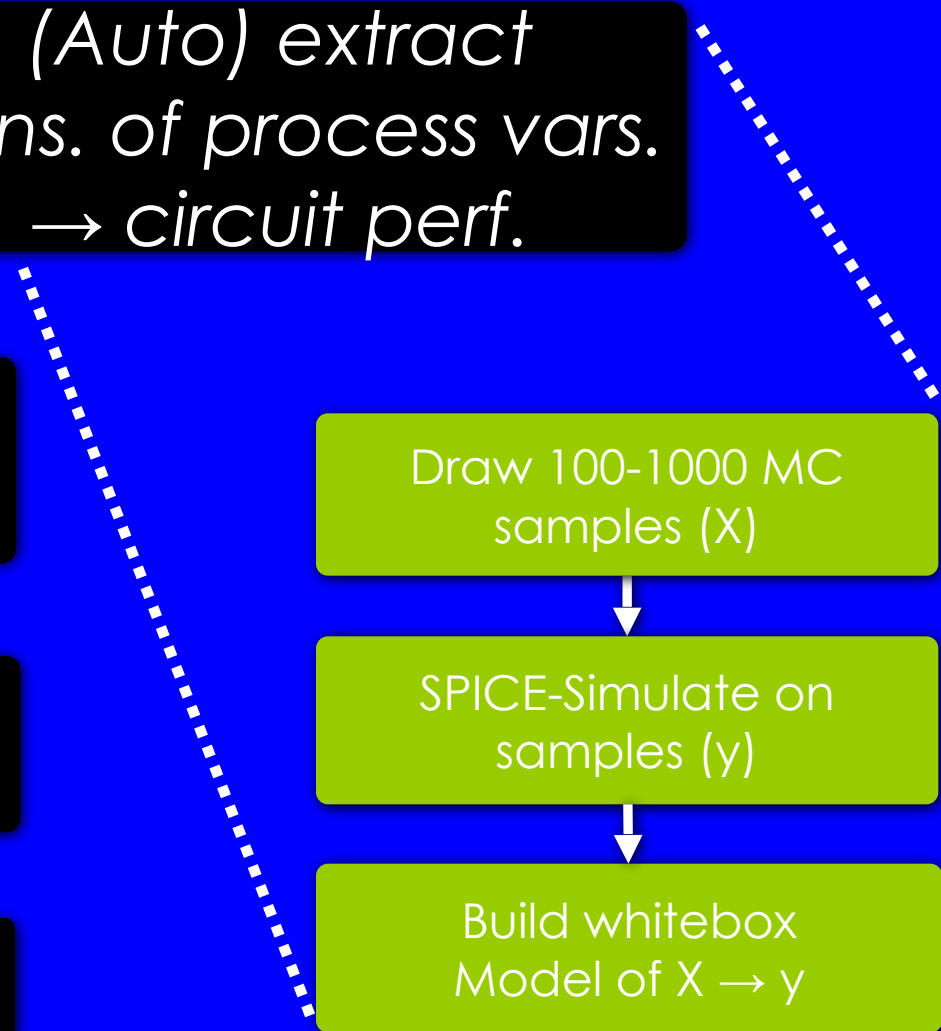


Layout

⋮

Draw 100-1000 MC
samples (X)


SPICE-Simulate on
samples (y)

Build whitebox
Model of $X \rightarrow y$




A New Modeling Algorithm: Fast Function Extraction (FFX)

1. Replace linear bases with a *ridiculous* number of nonlinear ones
 - x_i^2 , $x_i \cdot x_j$, \log , abs , $\max(0, x - \text{thr})$, ...
2. Pathwise regularized learn on bases



3. Filter for tradeoff of error vs.
coefficients



(That's it!) (+ a couple tweaks)

Outline

- Motivation
- Proposed flow
- Backgrounder
- Fast Function Extraction (FFX)
- Results
- Conclusion


Circuit Test Problems

Circuit	# Devices	# Process Vars.	Outputs Modeled
Opamp	30	215	AV (gain), BW (bandwidth), PM (phase margin), SR (slew rate)
Bitcell	6	30	cell _i (read current)
Sense amp	12	125	delay, pwr (power)
Voltage ref.	11	105	DVREF (difference in voltage), PWR (power)
GMC filter	140	1468	ATTEN (attenuation), IL
Comp-arator	62	639	BW (bandwidth)

Hspice™, industrial MOS models < 65nm, 800-5K MC samples

Circuit Test Problems


Circuit	# Devices	# Process Vars.	Outputs Modeled
GMC filter	140	1468	ATTEN (attenuation), IL



Replace linear bases with
a *ridiculous* number of nonlinear ones

x_i^2 , $\log(x_i)$, $\text{abs}(x_i)$, $\max(0, x - \text{thr})$,

$x_i \cdot x_j$, $\log(x_i) \cdot \text{abs}(x_j)$, ...



$$\approx (10 * 1468)^2 / 2$$

≈100M bases to select from (!)

Circuit Test Problems

Circuit	# Devices	# Process Vars.	Outputs Modeled
Opamp	30	215	AV (gain), BW (bandwidth), PM (phase margin), SR (slew rate)
Bitcell	6	30	cell _i (read current)
Sense amp	12	125	delay, pwr (power)
Voltage ref.	11	105	DVREF (difference in voltage), PWR (power)
GMC filter	140	1468	ATTEN (attenuation), IL
Comp-arator	62	639	BW (bandwidth)

Results summary: <30 s build time. Error always < linear & quadratic, sometimes dramatically.

Opamp PM Equations


(215 global + local process variables. Modeling time < 30 s)

# bases	Test error	Extracted Equation
0	15.5 %	59.6
1	6.8	$59.6 - 0.303 \cdot dxl$
2	6.6	$59.6 - 0.308 \cdot dxl - 0.00460 \cdot cgop$
3	5.4	$59.6 - 0.332 \cdot dxl - 0.0268 \cdot cgop + 0.0215 \cdot dvthn$
4	4.2	$59.6 - 0.353 \cdot dxl - 0.0457 \cdot cgop + 0.0211 \cdot dvthn - 0.0211 \cdot dvthp$
5	4.1	$59.6 - 0.354 \cdot dxl - 0.0460 \cdot cgop + 0.0198 \cdot dvthn - 0.0217 \cdot dvthp + 0.0135 \cdot \text{abs}(dvthn) \cdot dvthn$
...	...	
46	1.0	$58.9 - 0.136 \cdot dxl + 0.0299 \cdot dvthn - 0.0194 \dots$

Opamp PM Equations

Visualize Accuracy-Complexity Tradeoff

# bases	Test error
0	15.5 %
1	6.8
2	6.6
3	5.4
4	4.2
5	4.1
...	...
46	1.0



Voltage Ref. DVREF Equations

(105 global + local process variables. Modeling time < 30 s)

# bases	Test error	Extracted Equation
0	2.6 %	512.7
1	2.1	$504 / (1.0 + 0.121 \cdot \max(0, dvthn + 0.875))$
...
8	0.9	$476 / (1.0 + 0.105 \cdot \max(0, dvthn + 1.61) - 0.0397 \cdot \max(0, -1.64 - dvthp) + \dots)$

Shows: FFX is highly nonlinear if needed

Global vs. Local Variation?

FFX uses whatever variables help most, and sometimes patterns emerge

# bases	Test error	Opamp PM Extracted Equation
3	5.4	$59.6 - 0.332 \cdot dxl - 0.0268 \cdot cgop + 0.0215 \cdot dvthn$ global process variables

# bases	Test error	Comparator BW Extracted Equation
3	18.1	$171e6 - 4.57e5 \cdot x_{cm1,m1,lint} \cdot x_{cm1,m2,lint} + 5.23e4 \cdot x_{cm1,m1,lint}^2 + 4.80e4 \cdot x_{cm1,m2,lint}^2$ local (mismatch) variables

Highest-Impact Variables for Opamp PM

% Impact	Variable Name
46.5	dxl
10.2	cgop
9.7	dvthn
7.4	dvthp
3.9	RCN_nsmm_DXL
3.8	RCP_nsmm_DXL
3.6	dxw
3.1	cgon
2.3	RCP_nsmm_DXW
...	...
0.3	CM1_M1_nsmm_LINT
0.3	CMB2_M1_nsmm_NSUB
...	...

Outline

- Motivation
- Proposed flow
- Backgrounder
- Fast Function Extraction (FFX)
- Results
- Conclusion

Conclusion

- **Process variation is bad, and getting worse**
- There are solutions for design tuning
... But not for early-stage manual topology design
- **Idea: complement hand-derived equations with auto-extracted equations of variation → performance**
- **FFX builds the models (fast, scalable, nonlinear)**
...with the help of pathwise regularized learning
- Easy to get started: **code at trent.st/FFX**
 - Just ≈ 3 pages of python!
 - Solido using it extensively. Others too, for analog test, behavioral modeling, and even web search (!)